


Yet Another Introduction to Quantum Computing

Thomas R Clarke

September 19, 2026



Preface

In science fiction *quantum* is often used to explain away a technology that seems to defy the laws of physics. In the 2018 superhero movie *Ant-Man*, the protagonist is able to become smaller and smaller until they enter the mysterious *Quantum Realm*, described as “a reality where all concepts of time and space become irrelevant”. Perhaps the writers have a consistent set of logical explanations to describe how their worlds work. Our notions of space and time in our world might not be true in their world.

The set of rules that define quantum mechanics are a subject of intense creative destruction. Theoretically, defining space and time on the quantum level is a subject for higher dimensions and exotic symmetries. Experimentally, pushing the limits of technology to define every greater precision in measurement has given birth to the current information revolution. This revolution, like the industrial revolution, or the digital revolution that brought forth the world wide web, is as exciting as it is poorly understood. The quantum revolution offers exciting new technologies that could transform how we view the world (sensing), process it (computing) and share those insights (communications).

No single book could do justice to all the categories of quantum technologies, so this book will focus on computing in the most concise explanation, accepting the simple reality that the field is changing too rapidly for most of the content to be anywhere close to state-of-the-art. If all this book can accomplish for the reader is to develop an honest and intuitive appreciation for quantum information processing, it would have been worth writing. Least of all for myself.

Contents

1	How many physicists does it take to change a light bulb?	7
1.1	What happened to Schrödinger's cat?	8
1.2	Spooky action at a distance	8
1.3	Quantum Technologies	9
1.4	Chapter Summary	9
2	What is Quantum Computing?	11
2.1	More than flipping coins	11
2.2	Quantum Algorithms	11
2.3	Applications of Quantum Computing	12
2.4	Quantum Advantage	13
2.5	Chapter Summary	14
3	The Mathematical Minimum	15
3.1	Real and imaginary numbers	15
3.1.1	Complex numbers	15
3.1.2	The Argand Diagram	15
3.1.3	Complex conjugate	16
3.1.4	Magnitude of complex numbers	16
3.1.5	Interlude on Trigonometry	16
3.1.6	Euler's formula	17
3.1.7	Polar form of complex numbers	17
3.2	Linear Algebra	18
3.2.1	Vectors	18
3.2.2	Adding vectors	19
3.2.3	Multiplying a vector by a scalar	19
3.2.4	The inner product	19
3.2.5	Perpendicular vectors	20
3.3	Matrices	20
3.3.1	Matrix-vector product	20
3.3.2	Multiple matrices	21
3.3.3	Inverse matrices	21
3.3.4	Unitary matrices	22
3.4	Chapter Exercises	22
3.5	Chapter Summary	23
4	Dirac Notation	25
4.1	How to describe qubits	25
4.1.1	N-level quantum system	26
4.1.2	Bra: Row vectors with a complex twist	26
4.2	Measurement outcomes	27
4.3	Operators: A trip to the casino	28
4.4	Classical probability: What do we expect to get?	28
4.5	A quantum description of probability	28
4.6	Chapter Exercises	29
4.7	Chapter 4 Summary	29

5	Single Qubits	31
5.1	The unit of information	31
5.2	The Bloch Sphere	31
5.3	It's not just a phase!	32
5.4	Qubit gates	33
5.5	Parametrised quantum gates	33
5.6	Compiling any single qubit gate	34
5.7	Expectation values revisited	34
5.8	Gates vs operators, computation and physics	34
5.9	Bonus: Beyond Bits	34
5.10	Chapter Exercises	35
5.11	Chapter 5 Summary	35
5.12	References	35
6	Multiple Qubits	37
6.1	More than the sum of its parts	37
6.2	Reading binary	37
6.3	Multiple qubits, multiple gates	38
6.4	Product States	38
6.5	The Tensor Product	38
6.6	Entanglement	39
6.6.1	If statements with spicy sauce	39
6.6.2	Can the qubit not?	39
6.6.3	Ring the Bell (states)	40
6.6.4	General controlled unitaries	40
6.6.5	Multiple-control qubit gates	40
6.7	Chapter Exercises	40
6.8	Chapter Summary	41
6.9	Additional reading	41
7	Quantum Circuits	43
7.1	Introducing the circuit model	43
7.2	Basic building blocks	43
7.3	Quantum Random Number Generators	43
7.3.1	Why do we want random numbers?	43
7.3.2	The limitations of pseudorandom number generators	44
7.3.3	Quantum RNG	44
7.4	$1 + 1 = 2$ (most of the time)	44
7.5	Quantum Teleportation	45
7.6	Distributed Quantum Computing	45
7.7	Chapter 7 Summary	45
8	The Fourier Transform	47
8.1	Predicting the tides: From Newton to D-Day	47
8.2	The Fast Fourier Transform	48
8.3	A quantum mechanical treatment	48
8.4	A new limit of machine precision: Quantum Phase Estimation	49
8.5	Faster QFTs	50
8.6	Additional resources	50

Chapter 1

How many physicists does it take to change a light bulb?

In the early 20th century, a German physicist (by the cool name of Max Planck) was tasked with modelling the emission of light from a filament bulb. The bulb lights up when a wire gets hot enough to glow. This relationship between how hot the wire is and the colour of the light emitted is known as the black body radiation spectrum. It turns out, from his model, the energy emitted by an object with a finite temperature is not emitted in a continuous spectrum. That's to say rather than all possible wavelengths of light being emitted by this bulb only certain wavelengths were allowed. This is known as a discrete spectrum. Since each wavelength has a specific energy associated with it, as a consequence energy is emitted in chunks, or *quanta*. We can't have energy between these chunks. It's a bit like taking the lift (elevator in American English). One can only get on or off the lift at specific heights corresponding to the floor level. Each floor is separated by a distance and the lift can go up or down in any multiple of these distances limited by the number of floors. In between those levels the doors of the lift won't open.




Figure 1.1: Energy levels are quantised, like floors in a building, the lift can only exit on discrete levels.




Figure 1.2: As Max Planck discovered light emitted by a bulb is quantised, corresponding to specific energy transitions.

The quantisation of energy is where the word quantum comes from. Whilst this concept is less significant to quantum computing, it is still worth mentioning. The reason for this is that the transistor that powers all digital electronics is built upon this concept. Thus all computations using this technology (the integrated circuit) are computers running on quantum mechanics. This is not to say they are quantum computers simply that they are fundamentally built upon quantum physics, as are you, I, and everything we can see and touch. Quantum is nothing more than a limit, beyond which classical physics fails to provide an explanation. *Quantum technologies supersede the limits of what is possible with classical technologies.*

Beyond the quantisation of energy, there are quantum phenomena which are much more significant to quantum computing: *superposition & entanglement.*

1.1 What happened to Schrödinger's cat?

You may have heard of a very special cat associated with another German physicist, Schrödinger. In this analogy, a cat is kept inside a black box that the observer can't view inside. Within the box, there is also a vial of poison with a radioactive source. The radioactive source has an equal probability of decaying or not decaying in some time. If the source decays, the poison is released, and the cat dies. otherwise, the cat is just fine. Until the observer opens the box they don't know if the cat is dead or alive.




Figure 1.3: Schrödinger's cat thought experiment: the cat is in a superposition of alive and dead until observed.

[2]

Schrödinger came up with this analogy to explain why quantum mechanics was so controversial for physicists that had been able to do very well with classical physics until the turn of the 20th century. In quantum mechanics we call this particular state a *superposition*. And this is the state we say the Schrödinger's cat is in before the box is opened. This analogy is effective but is far removed from the everyday experiences one has, so another analogy can be made that encapsulates the same physics.

It is as simple as tossing a coin.

While the coin spins in the air, someone looking at it can't resolve which side is facing upwards. It is a blur of motion where the side facing up can't be determined. It is as if the coin were in a combination of being *both heads and tails at the same time*. This coin can be described as being in a superposition of heads and tails. In reality, there are not two coins, one with the heads facing up and another with the tails facing up, but only one coin with this special state. When the coin lands, it faces either heads or tails up. Now there is no uncertainty in the state of the coin. If we were to cover the coin and then look at it again sometime later, it would still have the same side facing up. The superposition the coin was in before has been collapsed. The concept of looking at the coin, or opening the box, is what in quantum mechanics is known as *performing a measurement*.

1.2 Spooky action at a distance

Historically, entanglement was described by Einstein as "spooky action at a distance". In a famous paper published in 1935, Einstein along with his coauthors, Boris Podolsky and Nathan Rosen, concluded that a quantum mechanical description of reality must be incomplete.

Instead of going through the theoretical and experimental physics developments, a simple analogy can serve as an introduction. Once the tools to describe qubits have been introduced, entanglement will be described in much more detail. For now, we continue with the coin analogy. The coins can work as a gross simplification, as the classical analogy can't accurately represent the intricacies of quantum theory. For this example, 2 coins will be required.

Tossing two unbiased coins is effectively the same as tossing the same coin twice and recording the outcome each time. Guessing whether one coin is H or T has a 50% success rate. To guess both coins correctly will happen 25% of the guesses. For the next step some removable adhesive, like Blu-tac, will be required. Imagine sticking the two coins together with each coin having the H facing outwards and the T's stuck together to make one coin twice as thick. Tossing

this combined coin will always result in the top coin having the H facing upwards. The coin at the bottom must also have the H side facing into the ground.



Figure 1.4: Two coins stuck together and spun can be used to describe entanglement.

Instead of revealing the combined coin, the two coins are separated. Since the coin at the bottom has H facing into the ground, it must have the T side facing upwards. This leaves us 2 coins one H facing upwards and the other T facing upwards. Without revealing the coins they are shuffled.

At this point, guessing either of the coins would have a 50% chance of getting the right outcome- the same as for one coin from a pair of simultaneously tossed coins. At this point, if only one of the coins is revealed, say the H coin, it is instantly known that the other one must be a T. Suddenly guessing the state of the second coin has a 100% chance of success whereas with the separate coin tosses revealing one made no difference to guessing the other. These linked probabilities are analogous (with some significant caveats) to an entangled quantum state. Entanglement can be thought of as these linked probabilities- it is impossible to describe one coin without describing the other.

Einstein argued that before the coins are revealed, they must have some side facing up. Looking at the coin only revealed information to us, the observer. In his view, some "hidden variables" were needed to make quantum mechanics fit with the classical intuition. Subsequent experiments demonstrated that this was not the case. It is as though both coins are not flat on the table but spinning in mid-air. Their spins are perfectly out of phase such that each coin is random right up until we see it.

1.3 Quantum Technologies

With some creativity, applications can be developed that offer some benefit to employing quantum effects. From the 1960's, first generation quantum technologies had a transformative impact on human development. The transistor, the development of Light Amplification by Stimulated Emission of Radiation, better known as LASER, and nuclear Magnetic Resonance Imaging (MRI) have reshaped our lives.

Where the second generation of quantum technologies come in is that they take advantage of superposition and entanglement. This book is almost exclusively focussed on using the principles of quantum theory for developing computational approaches. Quantum computing can be thought of as the quest to improve our existing problem-solving machines by harnessing these quirks of nature. Alternative second generation quantum technologies leverage superposition and entanglement to build more sensitive sensors and to send information with better security guarantees, just as the first generation did.

1.4 Chapter Summary

- The word quantum comes from quantised or in discrete units
- A superposition occurs when a system must be described in terms of all the possible states it can occupy
- Entanglement happens when the outcomes of two individually random systems are correlated more strongly than can be explained with classical physics
- Quantum technologies use superposition & entanglement in their operation

Chapter 2

What is Quantum Computing?

Before we add quantum spice to our computers, it would be wise to first describe what it is that our classical computers do. Computers work by executing logical operations (think addition/subtraction/multiplication/...) that take the machine from a starting point to a finishing point. The starting point that is fed into the computer is conventionally known as the input and what the computer ends up with is usually referred to as the output. An *algorithm* is a set of instructions that allow the computer to take the input and generate an output. For example, an algorithm for doubling a number could take an input of 3, multiply it by 2 and return an output of 6.

Going back to the coins here, this is a bit like laying them out in a row and then proceeding to shuffle and flip the coins according to a script that allows something to be done. For example, you could add 2 numbers expressed in binary by following a simple series of rules that allows for addition. This is similar to using an abacus as a calculator. Since the advent of the digital computer in the 1940's, these operations have been carried out by increasingly advanced machines to do more complex computations a lot faster.

2.1 More than flipping coins

In the last chapter, the analogy of flipping coins was used to illustrate the phenomena on which classical information and quantum information behave differently. Alas, our quantum computers of today cannot be built from a handful of loose change. They are generally quite expensive to buy. Before outlining how quantum computers execute quantum algorithms, let's first consider some drawbacks of our coin quantum computers:

- Flipping coins is not very fast. Modern computers have clock speeds in the GHz or billions of complete clock cycles per second. We would need an astronomical number of coins to try and replicate that.
- Tossing coins can only really be used to generate random numbers. Beyond the conventional deterministic computing operations we can already do, the only real advantage from quantum theory would be the 'random' nature of the coin flips.
- Sticking coins together is even slower than just flipping them. Not to mention the part where they get shuffled in a pseudo random way.
- Coins take up a lot of space. Imagine a piece of paper the size of a laptop. A few dozen coins could realistically be placed on that surface. For a few hundred coins, you'd need a small table. By comparison, flash storage can store billions or trillions of bits of information in a form factor smaller than a fingernail.

These coins represent the qubits in a quantum computer. We should be able to control these qubits faster and more precisely than with coins:

- Perform operations on qubits quickly
- Have a precise control over the probabilities
- Be able to effectively control multiple qubits at once
- Cause interference effects between the coins (this requires phase, to be introduced in chapter 5.1).

2.2 Quantum Algorithms

Conventional computers are really great at what they do. The field of high-performance computing has had many decades to mature. According to Top500, an index tracking the most powerful supercomputers, the most powerful

supercomputer, Fugaku is a \$1bn powerhouse. It consumes 29,899 kilowatts to power 7,630,848 cores. For context, a mid-range laptop today might have 6 cores and consume a total of around 20 watts of power. Both the laptop and the supercomputer execute algorithms that take some inputs and generate outputs. Whilst they are both excellent at solving a great many problems, there are two things they can't handle so well: superposition & entanglement.

Starting with a simple coin toss there are 2 possible outcomes. If we add a second coin there are 4 possible outcomes (HH, HT, TH, TT). As more and more coins are added, the number of possible outcomes doubles with each coin that is added. This number N , grows as $N = 2^n$. If we had 300 coins, that would mean a single toss has more possible outcomes than we estimate the total number of atoms in the universe. That's roughly 2 followed by 90 0's different combinations of H & T. Now imagine tossing these coins 300 times...

A quantum computer could handle this problem easily. Instead of needing a universe of atoms, 300 qubits of *sufficiently high quality* would be enough to run a quantum algorithm that could simulate these coins (including sticking of them together) to arbitrary precision. Quantum algorithms are similar to the input and output of classical algorithms but with the addition of superposition and entanglement in the middle. If an algorithm doesn't feature superposition & entanglement, it would always be better to use the classical computer.

*Quantum computers are **not** faster versions of regular computers*

*Quantum computers are **not** used to run conventional algorithms, or do things that are already very well done on a classical computer*

For a quantum computer to be useful it has to be able to execute some algorithm that serves a purpose and take a time that is short enough for the solution to be valuable. Going beyond the limits of the coins, we can consider what an ideal quantum information processor should be capable of. Such an ideal quantum processor would be built of a collection of ideal quantum, bits, called qubits, along with the necessary classical information resources needed to precisely and reliably control them. Such a list might include these points:

- Perform quantum operations on individual qubits instantaneously with perfect precision
- Be capable of inputting and outputting classical and or quantum information
- Generate entanglement between any and all of the qubits instantaneously and with unlimited precision

If physics were perfect, perfect qubits would exist in nature and we could just pluck qubits out of free space like taking grapes from the vine. Alas, the real world is a chaotic mess of information constantly coming into and out of existence. Everything takes some time, and nothing can be repeated the same way every time.

Whilst physics is far from perfect, the limits of physics are quite generous, at least in theory. Engineering can push the limits of technology ever closer to the physical limits. Meanwhile physicists derive clever tricks to beat the standard quantum limit and demonstrate these feats of precision experimentally. Physics and engineering requires the people, the time, and ultimately the funding to go beyond what is possible with today's technology.

2.3 Applications of Quantum Computing

For quantum computing to be useful, there has to be a well-defined problem where using a quantum computer will help. Figuring out which problems can benefit from quantum computing, how to effectively use the quantum computer, and even just how to build and control such a machine remain very open questions. These questions pose a great practical challenge in making QC useful. Even still, in 2021 over \$3 billion was invested into quantum computing. Why is there so much hype around such an early-stage technology?

There are many industries expecting significant value creation from quantum computing:

- Pharmaceuticals: Develop better medicine by simulating drugs more quickly & accurately
- Supply chain & logistics: Quantum computing to optimise supply chains
- Finance: Portfolio optimisation & fraud detection by simulation of stochastic variables
- Environmental: Simulation of chemistry for batteries, solar cells & materials
- Science: Simulation of quantum systems to better understand the universe

Across each of these areas, and many that have yet to be discovered, established tech companies and start-ups are working to come up with new ideas on how QC can be used. Predicting when these solutions might become available, or how valuable they might be, is a guessing game for investors and governments. In addition to potential direct impacts, there are a lot of potential benefits to developing the technology to engineer qubits. Being able to measure qubits better can help develop quantum sensors with applications in healthcare & navigation.

2.4 Quantum Advantage

If our conventional computers perform so many tasks better, there must be some motivation for quantum computing being such an exciting area of research. The maths is beautiful, the physics is exotic, and the engineers will never run out of problems to solve. But the funding has to come from somewhere. Quantum computing is expected to be more useful than a glorified coin tossing simulator. Despite the limitations in state-of-the-art quantum computing machines, it is known of a few valuable applications where quantum algorithms could *theoretically outperform* the best available classical solution. Under this condition, we could define quantum advantage. **Quantum advantage occurs when adding a quantum computer gives significant benefit to the user compared to the best available purely classical alternative.** The choice of "significant benefit" is arbitrary and ultimately up to the people who use the computer. For example, a few different ways a quantum computer could add value:

- Speed: Solve a problem faster for better decision making
- Accuracy: Provide more accurate solutions that are more useful
- Efficiency: Consume less energy over the entire process and thus cost less to run
- New approaches: Enable entirely new computational models & solutions to solve problems that previously couldn't be solved

For any potential benefit, there are two methods by which we can claim a quantum advantage. The most common approach is theoretical. Problems that can be modelled mathematically in such a way that is convenient for a quantum computer can be targeted. The motivation for using the quantum computer provides inspiration for an algorithm. Theoretical models give us algorithms where the number of steps required to solve the problem can be predicted. Given the total length of this computation, assumptions can then be made about the hardware. For instance, we could specify how quickly a quantum computer can execute layers of a circuit and then infer how long it would take to finish the algorithm. Such an analysis is usually called a resource estimate, and can motivate roadmaps for hardware companies.

The comparison with classical computing adds another dimension to the quantum advantage. There will always be classical overheads on any quantum circuit, so a quantum circuit of the same number of operations as a classical circuit will always be slower. Theoretical analysis can account for this by predicting the total time to solution, accounting for some fixed overheads.

For a small problem, for example, finding the shortest path in a small village, the overhead of using a QC will likely never become viable for any advantage. Instead of considering fixed problem sizes (where the problem size is usually described by how many possible solutions exist) the theoretical analysis considers a range of problem sizes. As the number of possible solutions increases, the time taken for the computer, classical or quantum, to find a solution gets longer. If there is a speedup for a quantum algorithm, there is some problem size for which the quantum algorithm can solve it in fewer steps. Once the problem is bigger than this crossover point, the advantage of using the quantum computer grows.




Figure 2.1: By making assumptions about how quickly classical and quantum algorithms can solve problems, quantum advantage can be predicted theoretically

Since classical computing is so much more mature, the timescales to execute steps of an algorithm can be on the order of nanoseconds (up to billions of steps per second). Quantum computers of today cannot match this speed and there remains a long way to reach comparable maturity. Taking into account limitations of the hardware makes quantum advantage much more difficult to achieve. Furthermore, since state-of-the-art quantum computers have fewer than 1,000 qubits (as of 2022), the maximum problem size that can be encoded onto the quantum computer remains very small. If

the crossover size is beyond the limits of the quantum computer, there can be no quantum advantage (at least from speed).

In recent years, hardware providers have made a big emphasis of increasing the scale of their quantum computers, predominantly by raising qubit counts. As of the time of writing, there has been no practical demonstration of a problem with real world (outside physics) problem where using a quantum computer was better than using the best classical solution. As quantum computers become increasingly powerful it is hoped that this threshold will be crossed in the next few years...

2.5 Chapter Summary

- Quantum computers are more than just random objects that have superposition
- Algorithms are sets of instructions that we give a computer to process information
- Quantum algorithms are like conventional algorithms but they allow us to use superposition & entanglement
- For some applications, quantum algorithms have shown to be much better than classical algorithms
- When a quantum computer completes a task better than the best available classical solution, there is quantum advantage

Chapter 3

The Mathematical Minimum

Up till now, this book has described quantum computing qualitatively. An introduction to the introduction. To continue this journey, it is essential to give a mathematical description that illustrates the beauty of the symmetry and showcases the power of the theory. To begin with, we must go beyond the limits of reality.

3.1 Real and imaginary numbers

So far everything in this textbook has been described using real numbers- that is the set of numbers you can produce on most calculators. For quantum computing, we require going beyond to complex numbers.

There are two square roots of 4: $+2$ & -2 . For numbers that aren't integers there are algorithms we can run on our calculators to work out the roots to arbitrary precision. In physics there are many situations in which the only valid solution to an equation calls for the square root of a negative number. To make this work, it is necessary to first define the square root of -1 as $i = \sqrt{-1}$.

Using some basic algebra we can factorise any surd (square root number) into a product of two surds $\sqrt{ab} = \sqrt{a} \times \sqrt{b}$. This holds true for negative numbers allowing us to take the square root of any negative number by breaking it down into the square root of -1 and the positive magnitude of that number. This can be written as:

$$\sqrt{-a} = \sqrt{-1} \times \sqrt{a}$$

$$\sqrt{-a} = i\sqrt{a}$$

For example the square root of -4 can be written as $+2i$, $-2i$. These numbers that can be written as ai are known as imaginary numbers.

3.1.1 Complex numbers

Imaginary numbers work similarly to the real numbers everyone uses on a daily basis. One can divide, multiply, add & subtract them just as you can for real numbers. What this also means is that any of $\div$, $\times$, $+$, $-$ operations can be applied to a combination of both real and imaginary numbers. For instance, the positive root of -4 , $2i$ can be added to 3 as $3 + 2i$. Any such number that can be written as a sum of a real and imaginary number is called a *complex number*. More generally we can define this mathematically as:

$$z = a + ib$$

Where a and b are purely real numbers. The real part of z is known as $Re(z) = a$ and similarly the imaginary component of is $Im(z) = b$.

3.1.2 The Argand Diagram

Complex numbers can be represented on a graph by splitting them into their real & imaginary parts. In the Argand diagram the x axis represents the real part and the y axis represents the imaginary part.

If a were negative, the real part would be negative so the arrow would point to the left instead. If b were negative, the imaginary part would be negative, so the arrow would point downwards instead.

3.1.3 Complex conjugate

Every complex number has a real & imaginary part. By taking the imaginary part of the complex number and making it negative, we get the complex conjugate. For example if the complex number z is

$$z = x + iy$$

The complex conjugate would be z^* written as

$$z^* = x - iy$$

In the figure below the complex number z is shown along with its complex conjugate. In this diagram $\bar{z}$ is used rather than z^* but it means the same thing.




Figure 3.1: The complex number z and its complex conjugate $\bar{z}$. Note how the real part is the same, but the imaginary part is on the negative axis.

3.1.4 Magnitude of complex numbers

It can be very useful to describe how big a complex number is. Since it has a real component and an imaginary component, the magnitude of a complex number indicates how big it is across both these components.

The magnitude of a complex number z is denoted by two vertical bars as $|z|$. It is calculated the same was as the Euclidean distance as


$$|z| = \sqrt{x^2 + y^2}$$

With $z = x + yi$.

The magnitude of a complex number is important for figuring out the probability of getting an output from a quantum computation. This will be explored in the next chapter.

3.1.5 Interlude on Trigonometry

Trigonometry is defined as *the branch of mathematics dealing with the relations of the sides and angles of triangles and with the relevant functions of any angles* but for our use it is much better to consider trigonometry the study of rotations. In later chapters, quantum gates will be described as rotations. Consider a circle with radius 1. Any point on the circle can be reached by rotating on the surface by an angle. The figure below shows how we can reach any point.



The point P is reached by the angle θ . It has coordinates $x = \cos(\theta)$ and $y = \sin(\theta)$.

In chapter 5, quantum gates will be described in terms of rotations around a higher dimensional sphere.

3.1.6 Euler's formula

Euler's formula is very useful and can be stated as

$$e^{i\theta} = \cos(\theta) + i\sin(\theta)$$

Where θ is an angle in radians. The significance of this is that one can represent any complex number in this fashion simply by multiplying the left-hand side by some magnitude. For example, the complex number $z = \frac{1}{2}(1 + \sqrt{3}i)$

Can be written as

$$z = \cos(\pi/3) + i\sin(\pi/3)$$

With Euler's formula this becomes


$$z = e^{i\pi/3}$$

In the context of quantum information this object $e^{i\theta}$ is called a phase and is of exceptional importance. To explain why θ is so significant would require its own chapter. Instead it shall be incorporated into later chapters as needed.

3.1.7 Polar form of complex numbers

We can represent any complex number on the *Argand diagram* in terms of its real and imaginary parts. From trigonometry, we can also represent any point on a circle using an angle and a given radius.

Euler's formula allows us to put these two together, representing any complex number as a point on a circle.



There are some complex numbers that are too big to be written as $e^{i\theta}$. For instance, the complex number $z = 5 + 12i$ cannot be written as $\cos(\theta)$ because $\cos(\theta)$ only goes up to 1. To overcome this, we can multiply by some magnitude to allow for bigger numbers. By multiplying by 13 we can write

$$z = 13(\cos(1.176) + i\sin(1.176)) = 5 + 12i$$

$$z = 13e^{i1.176}$$

Where did the 13 come from?

The factor in front of the e is the magnitude of the complex number. This is the same magnitude described in 3.2.5. This is found by taking the norm of the real and complex part

$$13 = \sqrt{5^2 + 12^2}$$

Using Euler's formula, any complex number can be written in terms of its magnitude and an angle. This gives us

$$z = x + iy = re^{i\theta}$$

Where r is the magnitude of the complex number $r = |z| = \sqrt{x^2 + y^2}$ and θ is the angle made between the real and the imaginary plane

3.2 Linear Algebra

The second part of chapter 3 will cover linear algebra. Linear algebra is the framework for quantum computing, so it is very important to be familiar with. Like trigonometry, linear algebra underpins most of physics, engineering and computer science. Its importance is difficult to overstate. In 1939, Paul Dirac reformulated quantum mechanics using linear algebra [1]. Linear algebra is so important that this way of describing quantum mechanics is known as matrix mechanics. There are two key elements to linear algebra: vectors and matrices.

3.2.1 Vectors

Vectors are objects that have both direction & magnitude. Velocity is a vector because it has a magnitude, the speed at which the object travels with, and a direction, where it's moving towards. We will use a different notation to represent a vector with a straight line and an angled bracket as $|v\rangle$.

Imagine a car travelling diagonally across a road. At the same time, the car is travelling along the direction of the road as well as perpendicularly across the road. It can be said that the velocity of the car has a component pointing along the road and another component across the road. What makes vectors useful to work with is the fact that a vector has components. These are the individual magnitudes of the vector in each direction. We can encode the speed of the car along the road (v_{alon}) and the speed of the car across the road (v_{acro}) in a vector with two components:

$$|v\rangle = \begin{bmatrix} v_{alon} \\ v_{acro} \end{bmatrix}$$

Here the vector $|v\rangle$ has two components, v_{alon} & v_{acro} so the vector has a dimension of 2. In 3D space, any point can be described by 3 vectors. In general, a vector can be of any dimension, with any number of components.

We can write any n-dimensional vector as

$$|v\rangle = \begin{bmatrix} v_0 \\ v_1 \\ \vdots \\ v_{n-1} \end{bmatrix}$$

Where the subscript v_i indicates the i^{th} component of $|v\rangle$.

Important note: We start counting from 0 and end at $n - 1$. This convention of counting from 0 is widely used in QC and conveniently is also used in python too.

3.2.2 Adding vectors

It can be useful to add vectors. All this requires is adding the corresponding components of each vector. Adding two vectors, $|a\rangle$ & $|b\rangle$

$$|a\rangle + |b\rangle = \begin{bmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{bmatrix} + \begin{bmatrix} b_0 \\ b_1 \\ \vdots \\ b_{n-1} \end{bmatrix} = \begin{bmatrix} a_0 + b_0 \\ a_1 + b_1 \\ \vdots \\ a_{n-1} + b_{n-1} \end{bmatrix}$$

The same can be done in reverse, a vector can be split into different components. For example

$$|a\rangle = \begin{bmatrix} 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 3 \\ 0 \end{bmatrix} + \begin{bmatrix} 0 \\ 2 \end{bmatrix}$$

This can be useful for expressing a vector in terms of its components.

3.2.3 Multiplying a vector by a scalar

Vectors can be made smaller or bigger by enlarging them by some scale factor. To make a vector bigger by some scale factor, the vector is multiplied by a scalar. To scale the vector $|v\rangle$ by some scalar quantity, a , this would be done by multiplying each component of the vector by the scalar factor.

$$a|v\rangle = a \begin{bmatrix} v_0 \\ v_1 \\ \vdots \\ v_{n-1} \end{bmatrix} = \begin{bmatrix} a \times v_0 \\ a \times v_1 \\ \vdots \\ a \times v_{n-1} \end{bmatrix}$$

3.2.4 The inner product

It can be very useful to compare two vectors by measuring how aligned they are. Imagine two vectors pointing the same way. These two vectors would have a lot in common, and so would be perfectly aligned. The measure of how similar two vectors are is the inner product. The inner product can be calculated by multiplying each component of the vectors and adding them up.

For example take the vectors $|a\rangle$ & $|b\rangle$ as

$$|a\rangle = \begin{bmatrix} 1 \\ 2 \end{bmatrix}; |b\rangle = \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$


The inner product between them is $\langle a|b\rangle$. This can be computed by taking the column vector $|a\rangle$ and turning it into a row vector where every element of $|a\rangle$ has taken the complex conjugate.

$$\langle a|b\rangle = [1^* \quad 2^*] \begin{bmatrix} 3 \\ 1 \end{bmatrix}$$

Here the 1st component of $\langle u|$ is multiplied by the 1st component of $|v\rangle$ and the second component of $\langle u|$ is multiplied by the second component of $|v\rangle$.

$$= (1 \times 3) + (2 \times 1) = 3 + 2 = 5$$

This gives a scalar quantity from two vectors and is used to indicate how aligned they are. The figure below shows this, where the length of $|a\rangle$ shows the component of $|u\rangle$ in the direction of $|v\rangle$. This is then scaled by the length of $|v\rangle$ to give the inner product.



The inner product depends on the angle between the vectors. If the vectors are pointing the same way, it will be bigger, if the vectors are pointing in completely opposite directions, it will be smaller.

3.2.5 Perpendicular vectors

If two vectors have no components in common, their inner product will be zero. This can be shown by two vectors that have an angle of 90° between them. We call these vectors perpendicular. In quantum mechanics, the preferred term is **orthogonal** which means the same thing but applies to more than just vectors.

For example the vectors $|p\rangle = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$; $|q\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ have a inner product of 0.

$$\begin{bmatrix} 1^* & 0^* \end{bmatrix} \begin{bmatrix} 0 \\ 1 \end{bmatrix} = (1 \times 0) + (0 \times 1) = 0$$

3.3 Matrices

Another object, closely related to a vector, is a matrix. A matrix is just an array of numbers which you can perform a matrix product on. All the quantum gates will be represented as matrices, so understanding how they work will be essential for quantum computations. Matrices are usually denoted by capital letters, for example M

$$M = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$$

M has 4 elements: 1,2,3 & 4.

Each element in M has an index: a pair of numbers that gives the position in the matrix. This is a bit like how map coordinates indicate a position. Starting from the top left, the index increases from left to right or from top to bottom. For example, the bottom left element (3) of M would have the index M_{10} as the first row, zeroth column. The choice of starting from zero is made to be consistent with binary and is also useful for programming languages like python that are widely used in quantum computing.

M could also be written as a collection of indices

$$M = \begin{bmatrix} M_{00} & M_{01} \\ M_{10} & M_{11} \end{bmatrix}$$

3.3.1 Matrix-vector product

Matrices, or more generally 2D arrays of numbers, are useful for storing information. The pixels on a screen have colour intensity values that can be efficiently stored as a matrix where the number of rows and columns is equal to the horizontal and vertical pixels on a screen (1080 by 1920 is a popular choice). As well as storing information, matrices can be used to transform vectors. This is done using the matrix-vector product. The matrix M acting on the vector $|a\rangle$ is denoted as $M|a\rangle$. The matrix multiplication is done by multiplying the corresponding elements of M & $|a\rangle$. Starting from the first row

of M we multiply out the elements of M with the elements of $|a\rangle$ starting from the left and add them up. Let's call the transformed vector $|a'\rangle$. We can compute $|a'\rangle$ by doing the matrix-vector product of M on $|a\rangle$ as

$$|a'\rangle = M |a\rangle$$

Explicitly calculating this gives us

$$\begin{aligned} |a'\rangle &= \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} \\ &= \begin{bmatrix} (1 \times 1) + (2 \times 2) \\ (3 \times 1) + (4 \times 2) \end{bmatrix} \\ |a'\rangle &= \begin{bmatrix} 5 \\ 11 \end{bmatrix} \end{aligned}$$

Matrix multiplication follows some of the same relationships as multiplying numbers, but not all of them.

3.3.2 Multiple matrices

A matrix transforming a vector can be thought of as an instruction. The original vector represents some initial starting position (like a coordinate x, y) and the matrix maps that point onto a new point (x', y') . Matrices can be applied in a sequence. Starting from the vector $|a\rangle$, applying the matrix A followed by B would look like

$$BA |a\rangle$$

In linear algebra the order is read from right to left, whatever is the furthest to the right happens first and whatever is furthest to the left happens last. Matrix multiplication follows some of the same relationships as multiplying numbers, but not all of them. In regular multiplication, the order in which numbers are multiplied doesn't matter. $5 \times 3 = 3 \times 5 = 15$. For any numbers a, b , $a \times b = b \times a$. The same cannot be said for any two matrices, $AB \neq BA$. This relationship, whether the order of operations matters, is defined as whether A and B commute. Integer multiplication is commutative, matrix multiplication does not usually commute.

3.3.3 Inverse matrices

For many matrix-vector products, we can go back and get the original vector from the transformed vector. This means, if we know the transformed vector $|a'\rangle$ & the matrix M we can work out the original vector $|a\rangle$.

To get the original vector, we do the matrix-vector product, but with the transformed vector and the inverse of the matrix. This gives us

$$M^{-1} |a'\rangle = |a\rangle$$

By using our equation for $|v'\rangle = M |a\rangle$ from earlier, this gives us

$$M^{-1} M |a\rangle = |a\rangle$$

Notice how $|a\rangle$ is completely unchanged from applying M and then M^{-1} . The operation $M^{-1}M$ effectively does nothing. The same would be true if we applied M^{-1} first and then M as MM^{-1} . The general name for the "do nothing" matrix is the identity. The identity is

$$M^{-1}M = MM^{-1} = I$$

Before we had $|a'\rangle = M |a\rangle$. And now we would like to go backwards from $|a'\rangle$ to $|a\rangle$. This is done using the inverse of M : M^{-1}

$$M^{-1} |a'\rangle = |a\rangle$$

Notice how we can take the original equation and do the matrix multiplication by M^{-1} and get the same result

$$M^{-1} |a'\rangle = M^{-1} M |a\rangle = |a\rangle$$

M^{-1} is cancelling out M so that there is no effect on $|a\rangle$. The matrix that represents this "no effect" is called the identity I defined by

$$I |v\rangle = |v\rangle$$

As a matrix, the identity looks like

$$I = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

3.3.4 Unitary matrices

For all quantum gates, an important property they have is that they are unitary. Unitary matrices have an inverse equal to their conjugate transpose (the dagger at the top).

$$U^\dagger = U^{-1}$$

This means

$$U^\dagger U = U^{-1} U = I$$

The conjugate transpose is the same process as it is to turn a ket (column vector) into a bra (row vector). Take You swap the rows and the columns and then replace every imaginary number with the negative version.

All quantum gates are unitary matrices, it is very common to see the term unitary when talking about any quantum circuit.

3.4 Chapter Exercises

1. Consider the complex numbers

$$z = 4 + 5i; y = -3 + i; m = -2 - 2i$$

- i What is $Re(z)$?
- ii What is $Im(m)$?
- iii Sketch z, y, m on the Argand diagram

2. Consider the complex number $z = 3 + 5i$.

- i What is the complex conjugate of z , z^* ?
- ii What is the magnitude of z ?
- iii What is the magnitude of z^* ?

3. Compute the inner product for the vectors

$$|w\rangle = \frac{1}{\sqrt{2}} \begin{bmatrix} i \\ 1 \end{bmatrix}; |z\rangle = \frac{1}{\sqrt{3}} \begin{bmatrix} \sqrt{2} \\ i \end{bmatrix}$$

- i. $\langle w|z\rangle$
- ii. $\langle z|w\rangle$
- iii. Bonus question: Compare i. ii. . Is the result the same as from Exercise 3.2 iv? If not, what has changed?

4. Imagine the vector $|v\rangle$ given by

$$|v\rangle = \begin{bmatrix} 2 \\ 1 \end{bmatrix}$$

- i What is the 0^{th} component of $|v\rangle$? What about the 1^{st} component?
- ii Write $|v\rangle$ in the form

$$|v\rangle = a \begin{bmatrix} 1 \\ 0 \end{bmatrix} + b \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

- iii Write down the vector $5|v\rangle$ and describe what has changed to the vector

5. Let the vector $|q\rangle$ be given as

$$|q\rangle = \frac{1}{4} \begin{bmatrix} -3 \\ 4 \end{bmatrix}$$

- i Using $|v\rangle$ from the last exercise, compute $|y\rangle = |v\rangle + |q\rangle$
- ii Compute the inner product $\langle q|v\rangle$
- iii Compute the inner product $\langle v|q\rangle$
- iv Bonus question: Compare your answers for ii. & iii. Explain any differences, if any, or why there is no difference. Does this hold for all pairs of vectors?
- v Using the result from either ii. or iii. are the vectors $|v\rangle$ & $|q\rangle$ perpendicular?

3.5 Chapter Summary

- Vectors are a collection of numbers stored in a row or column
- Vectors can be added to each other or multiplied by a scalar
- The inner product of two vectors tells us how aligned they are
- Perpendicular (orthogonal) vectors have an inner product of 0
- Matrices are arrays of numbers that can act on vectors
- Applying a matrix to a vector turns the vector into another vector
- Matrices can have an inverse which does the reverse of the matrix
- Unitary matrices have an inverse equal to their conjugate transpose

Chapter 4

Dirac Notation

This chapter explains a bit of the framework for quantum computing using the maths from the last chapter. It's like learning the language of quantum computing, first the rules of the grammar were introduced in chapter 3, now that grammar is being applied to create sentences.

4.1 How to describe qubits

When using a computer every day, it's important to be able to extract useful information from the computer. For instance, reading emails requires being able to access the information in your inbox. Unlike reading emails, quantum computing works at the lowest level of the system. It would be like directly accessing the voltages across transistors that make up bits in a classical computer. Instead of working with bits, quantum computing works with qubits. To do anything at all with qubits, it's important to describe what state they're in. The state of a qubit is described by something known as a ket.

How do we describe the state of our qubits? It is easy to think this would require a very complicated mechanism. Thankfully, it is as simple as a column vector! The state of our quantum computer can be described with a **column vector known as a ket**. Because it is a state represented by a vector, we call this the **state vector**. Usually, the state of our quantum computer is given by the ket $|\psi\rangle$.

The coins still work as an example for the state vector. If the coin is heads up we can use the ket $|H\rangle$, if it's tails up it is $|T\rangle$. When the coin is spinning in the air, the superposition state (see section 1.3), this state can be represented by the state $|\psi\rangle$:

$$|\psi\rangle = \frac{1}{\sqrt{2}}(|H\rangle + |T\rangle)$$

Since the coin has a $\frac{1}{2}$ probability of being in either heads or tails, you might have expected to see $\frac{1}{2}$ instead of $\frac{1}{\sqrt{2}}$. The reason for this is that in quantum mechanics, a state is described with an amplitude (the number outside the ket) proportional to the square root of the probability.

We can write out our kets as the simplest column vectors

$$|H\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

Similarly, the tails side up can be represented by the ket

$$|T\rangle = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

This allows $|\psi\rangle$ to be expanded as

$$|\psi\rangle = \frac{1}{\sqrt{2}} \left(\begin{bmatrix} 0 \\ 1 \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix} \right)$$

By adding the corresponding elements of each vector, this can be simplified to

$$|\psi\rangle = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = |+\rangle$$

⚡ This superposition state has its own symbol, $|+\rangle$.

Instead, if there were a biased coin which had a probability of heads given by P , we would have a different state. The probability amplitude for H is $\sqrt{P}$. The probability of getting tails would be $P(T) = 1 - P(H)$ which gives $P(T) = 1 - P$. So our probability amplitude for T is $\sqrt{1 - P}$. The state of our coin is then

$$|\psi\rangle = \sqrt{P}|H\rangle + \sqrt{1-P}|T\rangle$$

We can write this as a column vector as

$$|\psi\rangle = \begin{bmatrix} \sqrt{1-P} \\ \sqrt{P} \end{bmatrix}$$

4.1.1 N-level quantum system

In general, our quantum computer can take a lot more than 2 possible outcomes. Thankfully we can just add a row to our column vector for every possible outcome. So we can describe our state in terms of each possible outcomes it can take. Our quantum computer can output any number up to some maximum number of outcomes. We can write the state of any quantum computer in terms of all the possible states (numbers) we can get from it

$$|\psi\rangle = \sum_{i=0}^{N-1} c_i |i\rangle$$

Where the index i indicates a possible state we can observe the quantum system in. There are N such possible states. Each state has a probability of being measured given by $|c_i|^2$ with c_i being the probability amplitude. Writing this as a state vector gives us

$$|\psi\rangle = \begin{bmatrix} c_0 \\ c_1 \\ c_2 \\ \vdots \\ c_{N-2} \\ c_{N-1} \end{bmatrix}$$

A little note here: to model the state of a quantum computer requires keeping track of N complex numbers. That doesn't seem too bad. But for n qubits we have $N = 2^n$ complex numbers to keep track of. Just 20 qubits would have more than a million complex amplitudes to keep track of! Adding 1 more qubit to get 21 qubits gives us more than 2 million complex amplitudes.

4.1.2 Bra: Row vectors with a complex twist

Another important quantity to define is the row vector version of the state vector. This is called a **bra** and is represented by $\langle\psi|$. The big difference between normal row vectors and bras is that we always take the complex conjugate for every element of the column vector (ket).

For instance, the state with ket

$$|\psi\rangle = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ i \end{bmatrix}$$

Would have a bra given by

$$\langle\psi| = (|\psi\rangle^*)^T$$

First, all the imaginary numbers in $|\psi\rangle$ get a minus sign in front

$$= \frac{1}{\sqrt{2}} \left(\begin{bmatrix} 1 \\ -i \end{bmatrix} \right)^T$$

Then we transpose, shown by the T , by swapping the rows with the columns. We're left with our bra

$$\langle\psi| = \frac{1}{\sqrt{2}} [1 \quad -i]$$

In the same way, we can get the coin bras. The $\langle H|$ is the same as the ket $|H\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ but a row vector where the imaginary part of the complex numbers are all multiplied by -1 . Since there are no imaginary numbers in $|H\rangle$, the bra $\langle H|$ is just the transpose of the column vector as

$$\langle H| = (|H\rangle^*)^T = [0 \quad 1]$$

Just as with the column vector, we can write a general expression for any bra

$$\langle\psi| = [c_0^* \quad c_1^* \quad c_2^* \quad \dots \quad c_{N-2}^* \quad c_{N-1}^*]$$

4.2 Measurement outcomes

One of the most important features of quantum computing is measurement. Whenever we do a quantum computation, we change the state of the qubits. The strange thing is that we never actually see the state itself!

In the example with the coins, we never see two coins with opposite sides up. The superposition can't be observed directly. Instead, we perform a measurement on the qubits to get some information out of them. The measurement has outcomes with different probabilities, described by the state vector. It's like the lift analogy in Chapter 1, the lift is never observed between two floors, we only see it at the floors. We also don't observe part of the lift on one floor and part of the lift on another floor. Something has gone terribly wrong if the lift is split across multiple floors!

For the coins, one may wish to know the probability of flipping the coin and getting heads or tails up. With Dirac notation, the probability of an outcome (heads up) can be calculated using the inner product (see section 3.2.4 for a refresher on the inner product). Using the coin example, the probability of the coin ending up in the heads state $|H\rangle$ can be calculated as

$$P(H) = |\langle H|\psi\rangle|^2$$

The state we want to measure $|H\rangle$ is always on the left (in a bra) and the state we have, ψ is always on the right in a ket.

First we do the inner product $\langle H|\psi\rangle$

$$\begin{aligned}\langle H|\psi\rangle &= \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} \\ &= \frac{1}{\sqrt{2}}(1 + 0) = \frac{1}{\sqrt{2}}\end{aligned}$$

Then we get the probability by taking the magnitude and squaring

$$P(H) = \left| \frac{1}{\sqrt{2}} \right|^2 = \frac{1}{2}$$

Since all the probabilities of the coin landing on $|H\rangle$ or $|T\rangle$ must add to 1, for a biased or unbiased coin, it must be that $P(H) + P(T) = 1$. Let the probability of getting heads be P . This gives us a state of

$$|\psi\rangle = \begin{bmatrix} \sqrt{P} \\ \sqrt{1-P} \end{bmatrix}$$

For any quantum state, the probability of measuring it to be in the state that it is must be 1. For example, if the coin were heads up, and nothing happened to the coin, the probability of measuring heads would be 1. The probability of getting a tails would be 0. This can be verified by working out the probability of measuring tails. The probability of measuring tails $P(T)$ is still given by the same equation

$$P(T) = |\langle T|\psi\rangle|^2$$

where the coin is in the state $|\psi\rangle = |H\rangle$.

$$\begin{aligned}& \left| \begin{bmatrix} 0 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \end{bmatrix} \right|^2 \\ &= |(0 \times 1) + (1 \times 0)|^2 = 0\end{aligned}$$

On the other hand, the probability of measuring a heads is 1. Whilst this result may seem obvious, it is worth emphasising that for any quantum state, the probability of measuring it in that state should be 1 (neglecting measurement error). This leaves us the important result

$$|\langle \psi|\psi\rangle|^2 = 1$$

Another way of thinking about this is that for any quantum system, the probabilities of finding it in each of its states must add up to 1.

$\langle \psi|\psi\rangle = 1$ **must hold for any normalised quantum state.**

This simplifies a lot of the maths.

4.3 Operators: A trip to the casino

An operator is a mathematical object that acts on a ket and returns a new ket. Operators are matrices that allow us to process quantum states. This is how we do quantum computing. Our operators are denoted by a capital letter, such as $\hat{O}$. They act on a quantum state $|\psi\rangle$ to return a new state $|\psi'\rangle$ as

$$\hat{O}|\psi\rangle = |\psi'\rangle$$

This is just a matrix-vector product.

In gambling, the outcome of some event results in the gambler winning or losing money. Imagine a very simple game where every time you roll a heads you win \$1, every time you roll a tails, you lose \$1.

We can define an operator, let's call it $\hat{Z}$ that encodes the winnings and loses for each coin toss.

If we roll a heads $|H\rangle$ we should get +1, so we can encode that as

$$\hat{Z}|H\rangle = +1|H\rangle$$

Where the coefficient +1 is the winnings and the final state $|H\rangle$ is what we end up with.

Similarly for tails we lose a dollar so

$$\hat{Z}|T\rangle = -1|T\rangle$$

Using the vector representation of $|H\rangle, |T\rangle$ we can write the matrix $\hat{Z}$

$$\hat{Z} = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

It turns out $\hat{Z}$ is an important quantum logic gate. Note: since the symbol Z has no other major uses, it's common to see the operator without the hat as just Z .

We can invent another operator to perform the coin flip, we will call $\hat{F}$ for now¹ can be represented as

$$\hat{F} = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

4.4 Classical probability: What do we expect to get?

In life there are many processes where we want to predict something. For instance, a farmer would want to know how good their next harvest would be, or a stockbroker would want to know whether or not a stock is going to go up or down in price. In those cases, the farmer is not interested in the probability of getting an exact crop yield, but wants to know what yield they should expect. Quantum mechanics allows us to calculate these expectation values. What would we expect to get from our coin gambling example?

We know that half the time we will get \$1, and half the time we will lose \$1.

So our expected winnings, $\langle W \rangle$ would be

$$\langle W \rangle = \frac{1}{2} \times 1 + \frac{1}{2} \times -1 = 0$$

We win as much as we lose, so we'd expect to get nothing.

More generally we can write the expectation value in terms of the value of each outcome and the probability

$$\langle W \rangle = \sum_{i=0}^{N-1} X_i \times P(i)$$

Where X_i is the outcome for event i with probability $P(i)$.

4.5 A quantum description of probability

So far this description of probability is entirely classical, even the use of angled brackets! A quantum description is more elegant. A perceptive reader might have noticed the similarities between the probability

$$\langle W \rangle = \sum_{i=0}^{N-1} X_i \times P(i)$$

¹We will meet this seemingly trivial flip operator in Chapter 8 in a much more powerful form.

And the state vector

$$|\psi\rangle = \sum_{i=0}^{N_1} c_i |i\rangle$$

We can even get $P(i) = |c_i|^2$ so perhaps we can express the expectation value in terms of the state vector and something that gives us the outcome for each state, X .

The previous example with the \$1 gain/loss coin can be repeated with some quantum sauce. Recall the operator Z that gave us the win/loss for heads and tails. We would like to work out the expected earnings. We know that we can get the amount earned X_i

$$X_i(|i\rangle) = Z |i\rangle$$

Where i can either be H or T . And we can get the probability of getting i from $\langle i|\psi\rangle$.

Adding up all the probabilities gives us the important result

$$\langle Z \rangle = \langle \psi | Z | \psi \rangle$$

We have defined the expectation value of the operator Z . Assuming our coin is unbiased, it has the general superposition state. We can then compute the expectation value as

$$\begin{aligned} & \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \frac{1}{\sqrt{2}} \begin{bmatrix} 1 \\ 1 \end{bmatrix} \\ &= \frac{1}{2} \begin{bmatrix} 1 & 1 \end{bmatrix} \begin{bmatrix} 1 \\ -1 \end{bmatrix} \\ &= \frac{1}{2}(1 - 1) = 0 \end{aligned}$$

4.6 Chapter Exercises

1. Imagine a biased coin which is three times more likely to land on tails than on heads. Write down the corresponding state vector for such a coin.
2. Using the definition of the measurement probability, calculate the probability of measuring heads and the probability of measuring tails. Verify that these add to 1.
3. Flip the coin. By applying the flip operator to a coin that is heads up, show that it gives us the superposition state. Hint: compute the matrix-vector product of $\hat{F} |H\rangle$. Bonus question: what would you get if you did the same with the coin starting tails up?
4. Expectation values. For this question, we're graduating from a coin to a dice. Let's call N the number on a dice. Work out the expectation value by working out the number on each dice multiplied by the probability of rolling that number.

4.7 Chapter 4 Summary

- Quantum states can be represented by column vectors
- We can separate out a quantum state into orthogonal components
- The number in front of the component is the probability amplitude
- The probability of measuring a quantum system in a given state is given by the square of the modulus of the inner product
- The inner product of a quantum state with itself always has magnitude 1
- Operators transform quantum states and are represented by matrices
- The expectation value for any operator, $\hat{O}$ is given by $\langle \hat{O} \rangle = \langle \psi | \hat{O} | \psi \rangle$

Chapter 5

Single Qubits

5.1 The unit of information

Almost all articles or news reports talking about quantum computing start off by explaining how qubits are different from classical bits. Vague statements about superposition and entanglement try to impress the reader with a general idea of how quantum computers could be useful. This chapter applies the maths from previous chapters to properly describe superposition. Entanglement requires a later chapter to be introduced mathematically. Earlier, the coin example was introduced with the heads & tails outcomes we can observe. A qubit has the corresponding measurement outcomes $|0\rangle$ & $|1\rangle$, we can just replace $|H\rangle$ with $|0\rangle$ and $|T\rangle$ with $|1\rangle$.

$$|0\rangle = \begin{bmatrix} 1 \\ 0 \end{bmatrix}; |1\rangle = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

These two states are called the computational basis. Just as with the coins, our qubits can be in a superposition of $|0\rangle$ & $|1\rangle$. We can have any probability of measuring $|0\rangle$ & $|1\rangle$, so long as all the probabilities are real positive numbers which add to 1.

5.2 The Bloch Sphere

So far the linear algebra used makes for an efficient computational tool to describe what quantum states are and how they evolve. We can model how the qubits change using vectors and calculate probabilities with the inner product. One limitation of this Dirac notation is that it is difficult to visualise quantum states. A column vector is hard to think of intuitively. Thankfully, everything that has been used so far can easily be represented on a diagram. A sphere with radius of 1 can represent the quantum state of any qubit. This sphere is known as the Bloch sphere. Below, the quantum state $|\psi\rangle$ is represented on the Bloch Sphere.




Figure 5.1: The Bloch Sphere: a geometric representation of the state space of a qubit.¹

At the top of the sphere is the state $|0\rangle$, at the bottom is the state $|1\rangle$. Our state $|\psi\rangle$ is somewhere in between so it is a superposition of $|0\rangle$ & $|1\rangle$. The closer $|\psi\rangle$ is to the top, the greater the probability of measuring $|0\rangle$, the same is true for

the bottom with $|1\rangle$. The angle θ between the $|0\rangle$ & $|\psi\rangle$ gives us these probabilities.

Notice how $|0\rangle$ and $|1\rangle$ are separated by $180^\circ(\pi)$ on the sphere, but they are perpendicular, so the angle between them is actually $90^\circ(\pi/2)$. The angles on the Bloch sphere are twice as large as the angles for the qubits.

The angle θ on the Bloch sphere gives the probabilities with angle $\theta/2$

$$P(|0\rangle) = |\langle 0|\psi\rangle|^2 = (\cos \theta/2)^2$$

Similarly for the probability of $|1\rangle$

$$P(|1\rangle) = |\langle 1|\psi\rangle|^2 = (\sin \theta/2)^2$$

Up and down, the z-axis is where the $|0\rangle$ & $|1\rangle$ are, for that reason the z-axis is known as the computational basis. In addition we have the x & y axis which correspond to different states.

The states located on the ends of the axes are known as eigenstates. For Z these are $|0\rangle$ & $|1\rangle$. The X axis has eigenstates of $|+\rangle$ and $|-\rangle$, and the Y axis has eigenstates $|+i\rangle$ and $|-i\rangle$.

In quantum mechanics, eigenstates are defined by a particular operator. When the operator acts on its eigenstate the same state is returned, multiplied by an eigenvalue. The most famous example is the Schrödinger equation,

$$\hat{H}|\psi\rangle = E\psi$$

Here the operator $\hat{H}$ is known as the Hamiltonian, which acts on the eigenstate $|\psi\rangle$ to return the energy E associated with being in that state. Coming back to our qubits, we can use Z as an example operator. If the qubit is prepared in the initial starting state $|0\rangle$, we can apply the gate Z to $|0\rangle$. Since $|0\rangle$ is an eigenstate of Z , the resulting quantum state is simply $|0\rangle$ multiplied by 1 (the corresponding eigenvalue). Instead, if we started with $|1\rangle$, we would return $-1|1\rangle$. The eigenvalue of -1 cannot be measured, any more than the eigenvalue of $+1$ can be.

5.3 It's not just a phase!

To show probabilities, a single angle is enough. So why do we need a sphere? There is another important parameter that describes a qubit: the phase. At last, our coins will not work to describe phase, there is no analogue in classical computing, nor in simple probability. For example, earlier the flip coin operator F had this form:

$$F = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

In quantum computing, this operator is known as a Hadamard gate with symbol H . If we apply H to $|0\rangle$ we get the superposition state

$$H|0\rangle = \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) = |+\rangle$$

Applying it to $|1\rangle$ gives us a different superposition state:

$$H|1\rangle = \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle) = |-\rangle$$

We call this state $|-\rangle$ because of the minus sign between the $|0\rangle$ & $|1\rangle$

The probability of getting $|0\rangle$ or $|1\rangle$ is the same for both. If we were to measure either state, we would get 0 half the time and 1 the other half the time. Since the measurement gives the same outcome you might think they are the same state. They are different states, and surprisingly there is an easy way for us to tell which one we have.

What happens if we apply H on these states?

$$\begin{aligned} H|+\rangle &= H \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) \\ &= \frac{1}{\sqrt{2}}(H|0\rangle + H|1\rangle) \\ &= \frac{1}{\sqrt{2}} \left(\frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) + \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle) \right) \end{aligned}$$

The $|1\rangle$'s cancel each other out and the $|0\rangle$ add together to give us

$$= |0\rangle$$

So if we apply H once to $|0\rangle$ we get $|+\rangle$ and then applying it again goes back to $|0\rangle$. This implies applying H twice does nothing. $H^2 = I$.

So if we ever have $|+\rangle$ we apply H and get $|0\rangle$ 100% of the time. If we get $|1\rangle$ we must have had another state to begin with. If we had $|-\rangle$ would measure 1 100% of the time. In this example, where is the phase?

The phase is contained in the + or - sign between the $|0\rangle$ & $|1\rangle$.

The phase can be more than just ± 1 , it can be any complex number with magnitude 1. In general, any quantum state can have a phase ϕ

$$\alpha |0\rangle + \beta e^{i\phi} |1\rangle$$

With $0 \leq \phi < 2\pi$

We can use what we know from the Bloch Sphere to write α & β in terms of the angle θ

$$\cos(\theta/2) |0\rangle + \sin(\theta/2) e^{i\phi} |1\rangle$$

Remember Euler's formula (section 3.1.2) which allows us to convert from the exponential to the complex number

$$e^{i\phi} = \cos(\phi) + i\sin(\phi)$$

If we choose $\phi = 0$ we get $e^0 = 1$ which gives us

The phase is also shown on the Bloch sphere. In the above diagram the state $|\psi\rangle$ is described by two angles θ describing the probabilities of measuring 0 or 1 and ϕ describing the phase. The angle ϕ corresponds to a rotation by ϕ around the z-axis in the Bloch sphere.

Phases can be used to encode information, such as with the quantum Fourier transform, or to mark out quantum states (as with Grover's algorithm).

Phases only matter when they are relative.

5.4 Qubit gates

As mentioned in Chapter 4, operators are what allow us to do things to quantum states. Below are a collection of the most common single qubit operators.

Gate	Description	Effect on $ 0\rangle$	Effect on $ 1\rangle$	Matrix
I	Identity: do nothing	$ 0\rangle \rightarrow 0\rangle$	$ 1\rangle \rightarrow 1\rangle$	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
X	Not: Swaps $ 0\rangle$ with $ 1\rangle$	$ 0\rangle \rightarrow 1\rangle$	$ 1\rangle \rightarrow 0\rangle$	$\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$
Y	Y: Rotation by angle π around the y-axis of the Bloch sphere	$ 0\rangle \rightarrow i 1\rangle$	$ 1\rangle \rightarrow -i 0\rangle$	$\begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix}$
Z	Z: Adds a phase of π to $ 1\rangle$	$ 0\rangle \rightarrow 0\rangle$	$ 1\rangle \rightarrow - 1\rangle$	$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$
H	Hadamard: Creates the superposition state	$ 0\rangle \rightarrow \frac{1}{\sqrt{2}}(0\rangle + 1\rangle)$	$ 1\rangle \rightarrow \frac{1}{\sqrt{2}}(0\rangle - 1\rangle)$	$\frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$
S	Adds a phase of $\pi/2$ to $ 1\rangle$	$ 0\rangle \rightarrow 0\rangle$	$ 1\rangle \rightarrow e^{i\pi/2} 1\rangle$	$\begin{bmatrix} 1 & 0 \\ 0 & e^{i\pi/2} \end{bmatrix}$
T	Adds a phase of $\pi/4$ to $ 1\rangle$	$ 0\rangle \rightarrow 0\rangle$	$ 1\rangle \rightarrow e^{i\pi/4} 1\rangle$	$\begin{bmatrix} 1 & 0 \\ 0 & e^{i\pi/4} \end{bmatrix}$

Table 5.1: Common single-qubit gates and their action on the computational basis

Sometimes it is easier to perform gates by knowing what they do to the $|0\rangle$, $|1\rangle$ states, sometimes it is easier to do the matrix multiplication.

A more comprehensive list of gates can be found [here](#)

5.5 Parametrised quantum gates

In addition to these fixed quantum gates, there are quantum gates we can specify ourselves. These gates are described by a parameter we can control to change what the gate does. These gates are very similar to nodes in a neural network where we can change the weights (we will see a lot more about these in a chapter on quantum machine learning). For instance, we might want to get more than a 50% chance of measuring an output. We can see on the Bloch sphere that if we rotate our quantum state by some angle towards $|0\rangle$ or $|1\rangle$ we can boost the probability of measuring either state.

The matrix that does this rotation is known as R_y .

$$R_y(\theta) = \begin{bmatrix} \cos(\theta/2) & -\sin(\theta/2) \\ \sin(\theta/2) & \cos(\theta/2) \end{bmatrix}$$

The angle we give as input (on the Bloch Sphere) is θ the angle we perform the rotation is $\theta/2$.

We also have rotations for the x & y axis too. In fact, the X,Y gates are special cases of $R_x(\theta)$ & $R_y(\theta)$ with $\theta = \pi$. $R_z(\phi)$ (also known as P) is also useful because it allows us to give any phase to a qubit.

$$R_z(\phi) = P(\phi) = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\phi} \end{bmatrix} = \begin{bmatrix} e^{-i\phi/2} & 0 \\ 0 & e^{i\phi/2} \end{bmatrix}$$

This gives a phase shift of ϕ between $|0\rangle$ & $|1\rangle$.

5.6 Compiling any single qubit gate

Just as we can represent any single qubit state with two parameters, we can also represent any single qubit gate using parameters.

$$U(\theta, \phi, \lambda) = \begin{bmatrix} \cos(\theta/2) & -e^{i\lambda} \sin(\theta/2) \\ e^{i\phi} \sin(\theta/2) & e^{i(\phi+\lambda)} \cos(\theta/2) \end{bmatrix}$$

Just as we can reach any point on the Bloch Sphere with θ & ϕ we can get any quantum gate using U .

The set of all gates that can be made from U are known as the special unitary group $SU(2)$.

Since quantum computers are physical objects, not mathematical abstractions, U requires some mechanism by which it can be programmed into our qubit. In practice our $SU(2)$ matrices need to be mapped into a quantum control sequence defined by physical parameters. For qubits controlled by microwaves and/or lasers these physical parameters define the pulse that implements the operator on the qubit. To do this, an intermediate representation is defined allowing the user to input any quantum operator they want and then passing on to the control electronics a set of instructions.

Intermediate representations are effective because the linear algebra by which operators are defined allows U to be decomposed into a product of matrices. The matrices are chosen such that each matrix can be implemented effectively by the control electronic. We call these matrices native gates. Depending on how a qubit is built their native gate set will be different. This is especially true for entangling operations.

5.7 Expectation values revisited

In section ??

5.8 Gates vs operators, computation and physics

In the last chapter, operators were introduced as matrices that act on a quantum state. In this chapter, it seems quantum gates do the same thing. This is true, all quantum gates are operators, but not all operators are quantum gates. This book is written in the language of quantum information processing, where physical processes are abstracted away into idealised qubits. Physics is a theory of information, what it is that can be known about the universe, and so the language used, the mathematical framework, is generally equivalent to our idealised computations. As always, there are applications for which small details can become more convenient to adopt the physics perspective.

For instance, the creation operator $\hat{a}^\dagger$ creates a photon from the vacuum but can't be represented by a $SU(2)$ matrix. The momentum operator $\hat{p} = -i\hbar \frac{\partial}{\partial x}$ is a partial derivative existing in a continuous momentum Hilbert space of momentum states. Being able to map our physical world onto qubits, or moving from the physics perspective to the computational perspective, is the key skill in making use of quantum computers. Doing this mapping enables quantum simulation. It generally considered that quantum computers are more effective than classical computers in this mapping.

5.9 Bonus: Beyond Bits

For this course, we work with qubits- 2 level quantum systems. There is considerable research into using more than 2 levels for quantum computation. Such systems are called qudits: d-level quantum systems that we can use to process quantum information.

Qudits are much less commonly used than qubits for a few reasons

- Having multiple levels is much more complicated. Think about the X gate, for a qudit, this gate becomes much more complicated

- Qudits present even more engineering challenges than qubits. For instance, a qubit in a trapped ion consists of the ground level and the first excited state. Higher energy levels have smaller gaps in energy, so they are more difficult to control precisely. This also introduces more complex physical error models.
- Quantum error correction for qudits is much more complicated. With qubits $|0\rangle$ can only go to $|1\rangle$ but for a 3-level system there are 2 possible states $|0\rangle$ could go to. And that's just one qutrit (a qudit with $d = 3$). Correlations between all 3 states of two or more qudits means the different combinations of errors scale exponentially. This is also true with qubits, but at least they have a base of 2 rather than 3 or more.

5.10 Chapter Exercises

1. Where on the sphere? We know the superposition state $|+\rangle$ has equal probabilities of $|0\rangle$ & $|1\rangle$. Where on the Bloch sphere would this state be represented? What is the angle θ that would describe this state?
2. Back to the start. Apply H twice to $|1\rangle$ and verify the state we end up with is $|1\rangle$.
3. The coin doesn't work. When we apply the Hadamard twice to a state, we get the same state back. Without doing any maths, consider throwing a coin. Why does throwing the coin twice not guarantee you get back the side you started with?
4. By considering the effect on $|0\rangle$ & $|1\rangle$, or by doing matrix multiplication, calculate the following states:
 - a. $XZ|1\rangle$
 - b. $XH|+\rangle$
 - c. $XYZ|0\rangle$
5. By using $\theta = \pi/2$ apply $R_y(\pi/2)$ to the state $|1\rangle$. What gate does $R_y(\pi/2)$ do?
6. The X gate turns $|0\rangle$ to $|1\rangle$ and $|1\rangle$ to $|0\rangle$. By considering angles on the Bloch sphere, what is the angle θ such that $R_y(\theta) = Y$
7. Which gate in the list above is represented by $U(\theta, \phi, \lambda) = U(0, 0, 0)$

5.11 Chapter 5 Summary

- The state of a qubit can be represented on the Bloch sphere
- The BS shows the probability and phase of a qubit
- The phase of a qubit is the $e^{i\theta}$ difference between the $|0\rangle$ & $|1\rangle$ states
- There are many qubit gates we can use to perform operations on qubits
- Some qubit gates are rotations specified by a parameter which describes what the gate does

5.12 References

[1] Chetan Wake, 2019, Introduction to quantum computing part -1 Representation of qubit using Bloch sphere

Chapter 6

Multiple Qubits

6.1 More than the sum of its parts

If quantum computing was nothing more than a collection of qubits processed individually, there would be little point in building a quantum computer. A single qubit can be modelled using 2×2 matrices - you can even work these out by hand! Even if we were to have a large number of these in parallel, it would still be very easy to simulate a large number of qubits. You might prefer to use a classical computer that is very efficient at evaluating linear algebra rather than pen & paper.

What makes qubits difficult to simulate is the interactions between them. Having qubits interact with one another allows for more complex algorithms. Thanks to entanglement, one qubit can control what happens to another. This has some interesting implications for quantum algorithms.

6.2 Reading binary

To describe many qubits, we need to add something to Dirac notation. For instance, if we were to label our qubits alphabetically, we'd run out of letters after just 26 qubits. Thankfully, we can leverage the binary system from classical computing to describe as many qubits as we want to. Everyone is familiar with the decimal system where counting goes as 0,1,2,...,8,9 and then we repeat from 10,11,12,... . Computers, classical or quantum * count in base 2. Binary works the same way as the decimal system but instead the count repeats every 2. Numbers go as

0,1,10,11,100,101,110,111

Which in decimal would be

0,1,2,3,4,5,6,7

To read a binary number:

1. Count the number of digits in the number and take this number - 1 as n . 2. Starting from the first digit on the left, if that digit is 1, calculate 2^n . 3. Moving to the next digit on the right, if that digit has a 1, calculate 2^{n-1} . 4. Add the result of step 3 to step 2. 5. Repeat this process for each digit decreasing the power by 1 each time until the last digit (which should have a $2^0 = 1$ calculation). An example of this is shown below



1. Here there are 8 digits, so we take $n = 8 - 1 = 7$. 2. For the first digit there is a 0, so we do nothing. 3. The second digit is 1 so we do $2^{7-1} = 2^6 = 64$. 4. Adding $0 + 64 = 64$. 5. The sum we get is

$$0 + 2^6 + 0 + 2^4 + 2^3 + 0 + 0 + 2^0 = 89$$

Here we needed at least 6 digits in binary to represent a number with 2 digits in decimal. The decimal notation is much more efficient. Let's say we wanted to encode the number 5 onto a quantum state.

We know $5 = 2^2 + 2^0$ so we can write 5 as 101 in binary.

As there are 3 digits, we need 3 qubits to represent this number.

The qubit on the left is in state 1, the qubit in the middle is in state 0 and the qubit on the right is in state 1.

$$|101\rangle = |q_2 q_1 q_0\rangle$$

Here we use the convention where the first qubit (2^0) is on the right, with the next qubit one to the left and the last qubit the most to the right. This is what IBM uses and what we will use when programming quantum computers.

Quantum states are read right-to-left with the first qubit next to the

This right to left notation can also be represented as a ket as $|101\rangle_{2,1,0} = |1\rangle_2 |0\rangle_1 |1\rangle_0$ where the subscript index specifies which qubit is in which state. Similarly, we can have our quantum gates given these subscripts so we can say which qubits they are acting on.

6.3 Multiple qubits, multiple gates

If we have multiple qubits, we can apply single qubit gates on each of them at a time. This is the same as having a bunch of single qubits on their own, each having their state changed by the qubit gate. The simplest case is two qubits. Let's begin with our qubits in the normal starting state of $|q_1 q_0\rangle = |00\rangle$. We can apply a Hadamard gate on the first qubit, q_0 to change it to the superposition state $|+\rangle$

$$|00\rangle \rightarrow |0+\rangle$$

Then we might want to apply the not (X) gate to the other qubit, q_1 to change it to $|1\rangle$

$$|0+\rangle \rightarrow |1+\rangle$$

As these are two separate operations, we could do them both at the same time in one step

$$|00\rangle \rightarrow |1+\rangle$$

Applying single qubit gates to a collection of qubits, changes the states of the qubits the gates are applied to

6.4 Product States

So far we have acted on all our qubits separately, what happens to one doesn't affect the others. If we were to measure one of the qubits, it wouldn't tell us anything about the other qubits.

Such states where we can separate the qubits out are called product states (or pure states).

In general a product state can be written as

$$|\Psi\rangle = |\psi_0\rangle |\psi_1\rangle \dots |\psi_{n-1}\rangle$$

Let's say we measure the first qubit (qubit 0) and get a 0. The remaining qubits wouldn't be affected, so our new state would be

$$|\Psi\rangle = |0\rangle |\psi_1\rangle \dots |\psi_{n-1}\rangle$$

6.5 The Tensor Product

Mathematically there is a very satisfying way of representing this with linear algebra. The tensor product allows us to do this. At this point some subscripts can be useful for labelling our qubits. The first (or zeroth) qubit can be denoted by a small a on the right of the ket, the second (or first) qubit will have a little b on the right. If qubits a and b are not entangled, the system of ab can be described as a tensor product.

$$|\psi_a\rangle_a |\psi_b\rangle_b = |\psi_a\rangle_a \otimes |\psi_b\rangle_b$$

Where the symbol $\otimes$ represents the tensor product operation.

Expanding our kets into column vectors

$$|\psi_a\rangle_a = \begin{bmatrix} a_0 \\ a_1 \end{bmatrix}; |\psi_b\rangle_b = \begin{bmatrix} b_0 \\ b_1 \end{bmatrix}$$

$$|\psi_a\rangle_a \otimes |\psi_b\rangle_b = \begin{bmatrix} a_0 |\psi_b\rangle_b \\ a_1 |\psi_b\rangle_b \end{bmatrix}$$

This is a new vector with $|\psi_b\rangle_b$ repeated twice, each time multiplied by a scalar b_0 or b_1 . The final vector has 4 elements and looks like:

$$|\psi_a\rangle_a \otimes |\psi_b\rangle_b = \begin{bmatrix} a_0 b_0 \\ a_0 b_1 \\ a_1 b_0 \\ a_1 b_1 \end{bmatrix}$$

So you can think of your new state as a 4-level quantum system. Every time we add another qubit, we tensor product it in, so the size of the vector doubles. This is where the exponential scaling of the state-space comes from. The same applies for multiple quantum gates in parallel. Our matrix increases by a factor of 2 in both the number of rows or columns. This is even better scaling!

Keeping with the number 5, the state vector for this number

$$|5\rangle = |101\rangle = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

For each qubit we add, the number of rows in this column vector doubles. This is what makes it difficult to simulate quantum algorithms on classical computers.

6.6 Entanglement

You might have noticed that it's really easy to apply single qubit gates to a collection of qubits. In the last chapter, we saw that all single qubit gates can be represented by 2×2 matrices. And each qubit can be represented by a column vector with 2 rows. So if we had a large number of qubits, we could still simulate a large number of 2×2 matrices on a classical computer. If it is easy to simulate a collection of isolated qubits, why do we bother making quantum computers?

Alone superposition has quite limited applications. Applying single qubit operations can be used for sampling distributions, but this limits quantum computers to being random number generators. Where superposition can become more useful is when entanglement enters the picture.

6.6.1 If statements with spicy sauce

In general programming, if statements allow you to include logic in a program. If a statement is true, the algorithm will do one thing, otherwise, if the statement is false, it will do something else. The computer first has to check the variables value before executing the next step of the algorithm. We can do this on a quantum computer by measuring our qubits, collapsing the superposition and then performing our next quantum gates accordingly. For example, in quantum teleportation this is done.

Back in section there was a game where if a coin landed up we got +1 dollar, otherwise (if it were tails up) we got -1 dollars. This works as a classical if statement, we add the value after the coin is landed and observed. What makes quantum operators different is that we can apply them to the superposition. What this means is we can act on both the $|H\rangle$ and $|T\rangle$ amplitudes at once.

All the way back in Chapter 1, entanglement was introduced

In general English, we use the word entangled to suggest that a process has mixed things up in such a way that they are hard to separate. This definition works nicely for quantum states where we cannot describe one qubit without describing another. What this means, is that when two qubits are entangled, the superposition of one qubit defines the superposition of another. Mathematically, we can't treat our qubits as a product state of separate 2×1 vectors any more.

6.6.2 Can the qubit not?

Entanglement can be created by gates that act on multiple qubits together to create an entangled state. The most common way this is done is by using a 2 qubit gate called the controlled-not, usually shown as CNOT or CX. The CNOT acts on two qubits, the control qubit and the target qubit. If the control qubit is in the state $|0\rangle$, it will do nothing, if the control qubit is in the state $|1\rangle$ it will perform the not (X) gate on the target qubit. In both cases, the state of the control qubit is unchanged. We can represent this with a table

Control	Target	→	Control	Target
0	0		0	0
0	1		0	1
1	0		1	1
1	1		1	0

If we have two qubits $|\psi_i\rangle_b |\psi_i\rangle_a$ where qubit 0 is in the state $|\psi_i\rangle_a$ and qubit 1 is in the state $|\psi_i\rangle_b$, there are two CNOT's we can have. One where qubit 0 is the target and qubit 1 is control or the other way around, qubit 1 is the control and qubit 0 is the target. Each of these has a different 4×4 matrix

The first has the control on the

$$CNOT_{0,1} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix}$$

$$CNOT_{1,0} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

6.6.3 Ring the Bell (states)

There are many ways we can entangle two qubits, and some qubits can be more entangled than others. The most common entangled states are known as the Bell states, and can be written like


$$|\Phi^+\rangle = \frac{1}{\sqrt{2}} (|00\rangle + |11\rangle)$$

$$|\Phi^-\rangle = \frac{1}{\sqrt{2}} (|00\rangle - |11\rangle)$$

$$|\Psi^+\rangle = \frac{1}{\sqrt{2}} (|01\rangle + |10\rangle)$$

$$|\Psi^-\rangle = \frac{1}{\sqrt{2}} (|01\rangle - |10\rangle)$$

For each of these states, we can't separate out our two qubits as a product state. The following circuit generates the Bell states



6.6.4 General controlled unitaries

Instead of having a not-gate being applied, we can have any single quantum gate. The logic follows from the controlled not. For the control qubit in the state $|0\rangle$ no action is carried out, whilst the control qubit in the state $|1\rangle$ will implement the unitary on the target.

6.6.5 Multiple-control qubit gates

We can have controlled quantum gates where there are more than one qubit as control. In this case, the target qubit is only affected if all the control qubits are 1.

The Toffoli gate is the same as the controlled-not gate but with two qubits needing to be $|1\rangle$. Only if both the control qubits are in the state $|1\rangle$ for the X to apply on the target qubit. Otherwise the Toffoli gate acts as identity.

6.7 Chapter Exercises

1. What would the following numbers be in decimal?
 - a. 1111
 - b. 1001
 - c. 101101
 - d. 0110

2. What would the following numbers be in binary?
 - a. 1
 - b. 5
 - c. 6
3. For n qubits, we need a vector with 2^n rows to describe the state. What shape of matrix (rows, columns) do we need to describe the operator (i.e. multi-qubit gate) that acts on this vector?

6.8 Chapter Summary

- Quantum computers work in binary
- Many qubits have quantum states represented as a binary number
- Quantum states are read right-to-left with the first qubit at the end
- We can apply single qubit gates individually across multiple qubits
- When qubits are independent of one another, they are in a product state
- When one qubit influences the state of another qubit they are entangled
- We can use entanglement to apply logic gates

* There have been proposals for using a base other than 2 for quantum computing. This has some advantages but the difficulty in constructing algorithms combined with the difficulties in error correction such a device has prevented any significant machine from being built.

6.9 Additional reading

[1] L.Vivah, Learn How to Read Binary in 5 minutes, 2018

Chapter 7

Quantum Circuits

For this chapter, it is recommended to view the Jupyter notebook version as a supplementary resource. You can find the link [here](#).

7.1 Introducing the circuit model

So far the description of quantum computing has been very algebraic. States are represented as vectors, gates are matrices, and by applying the corresponding indices we can implement single qubit gates on an array of qubits. This description has been fine so far because linear algebra is very efficient when it comes to single qubit operations. Even a CNOT can be done quite easily. Where linear algebra becomes inconvenient is when we want to talk about how we can solve problems, i.e. useful computation. For this we can turn to a new model. Using diagrams to encode a model can be more intuitive, for example, an electronic circuit can make it easier to represent a large and complicated electronic system. In much the same way, a quantum circuit is a diagram showing the operations that act on qubits. Whereas electronic circuit diagrams act like a shorthand that can be useful for calculating currents and voltages, quantum circuits act like a shorthand to help keep track of the states of qubits. All quantum circuits have 3 parts:

1. Initialisation: Start your qubits in an initial quantum state (usually in $|0\rangle^{\otimes n}$)
2. Operation: Perform a series of single qubit gates and entangling operations.
3. Measurement: Measure the qubits and return their projected states as classical variables (a bitstring)

Since classical bits are an abstraction of quantum bits (in the limit of continuous measurement), we can generalise the aforementioned steps to any classical algorithm too. In practice, any quantum computer will need to interact with a classical computer, so the circuit model is also convenient in representing cases where classical bits can be interfaced with quantum bits.

7.2 Basic building blocks

Just how the Bloch sphere can be used to represent states and gates, so too can circuit diagrams be used to encode the steps for any quantum circuit.

7.3 Quantum Random Number Generators

Pick a random number between 1 and 10.

Was it 7?

Humans are really bad at picking random numbers. For instance, choosing 1 or 10 doesn't seem so random because they are the largest and smallest numbers. A number picked near the middle intuitively feels more random than one at the higher or lower end. Even numbers seem less random than odd ones (though there is no reason for this to be true).

7.3.1 Why do we want random numbers?

For much of human history, random numbers were only used in games of chance. Dice go back 5,000 years (Piovano, 2011/2016). During WW2, random numbers became an important statistical tool for von Neumann when he was working on the Manhattan project (Metropolis, 1987) and for use by the Germans in sending encrypted messages. From the Manhattan project (Metropolis, 1987) came Monte Carlo methods. These are simulations that take a large number of

samples from a model using random numbers to compute something that would be difficult to solve otherwise. They are a powerful and ubiquitous tool in physics, economics and data science. It became clear that random numbers were increasingly useful in areas of science, cryptography and statistics. More recently, with the abundance of private information sent over the internet, there is considerable need to generate a large number of random numbers for various encryption standards (Zhou & Tang, 2011).

7.3.2 The limitations of pseudorandom number generators

Pseudo-random number generators (PRNGs) do not actually produce random numbers. Anyone who knows the seed (initial random number) and the algorithm can generate the entire sequence with complete certainty. If left to run for long enough, both the middle square method and linear congruential method will repeat. This is fine for video games where it is the feeling of randomness which matters, but not so great for encrypting communications (Li, 2013). With insider knowledge of the PRNG, an attacker could decrypt the communication.

7.3.3 Quantum RNG

Thankfully there are ways of generating truly random numbers based on physical processes. Atmospheric noise, the cosmic microwave background (the effect that caused static on old TVs) and radioactive decay are good examples. We can measure radio wave or microwave radiation, or the number of clicks from a Geiger counter. Quantum computers can also be used for this purpose. They are effectively controlled physical experiments leveraging quantum mechanics to perform some computation. Since randomness is an inherent part of quantum mechanics, quantum computers, unlike classical ones used by von Neumann, can serve as a True Random Number Generator (TRNG) (Jacak, 2021). This is because a quantum system can exist in a superposition of possible states, and following a measurement takes on one of these states. Whilst we can know the probability of the system taking each of these states, we cannot know with absolute certainty which it will take (Nielsen & Chuang, 2000).

Below is an example of how n random numbers (from 0 to $n_{\max}$) can be generated using IBM's quantum computer. The code creates a quantum circuit with enough qubits to represent the power of 2 greater than $n_{\max}$. The qubits are then put into a superposition and measured to obtain the random number. This process is repeated 1000 times and sampled n times to produce the numbers.

The TRNG implemented here is hardly the most practical implementation. It is rather slow, requires access to IBM's cloud infrastructure, is vulnerable to interception, and runs on a device cooled to nearly absolute zero. A more useful implementation of this concept was done 20 years ago by the Swiss company ID Quantique, using a photonic chip. Newer models can be integrated in desktop PCs with PCIe connectivity or even USB (ID Quantique, n.d.).

7.4 $1 + 1 = 2$ (most of the time)

A simple quantum circuit can be used to add 2 bits. This example illustrates a basic quantum algorithm, but you would never want to do this. There is no advantage in doing addition on a quantum computer, but this is a useful example of a quantum circuit. The adder should do the following: take two qubits and return their sum. This is simply adding $1 + 1$ in binary.

Input 1	Input 0	→	Output 1	Output 0
0	0		0	0
0	1		0	1
1	0		0	1
1	1		1	0

How do we construct a circuit that can implement this table?

For both inputs and outputs, each bit of information is represented as a qubit. The states of the output qubits, q_2, q_3 must be determined by the states of the input qubits to be added q_0, q_1 . The circuit that implements this requires controlled nots and a new gate known as the Toffoli gate.

For the circuit below, we'll have the top two qubits be the input, and the bottom two qubits be the output. This means when we measure our qubits, the bits to the left will be the output and the bits to the right will be the input.

7.5 Quantum Teleportation

7.6 Distributed Quantum Computing

7.7 Chapter 7 Summary

- A quantum circuit is a diagram that encodes all the steps of a quantum algorithm
- Quantum algorithms require 3 steps:
 1. Initialisation: setting all the qubits to $|0\rangle$
 2. Gates: a series of single qubit and entangling (multi-qubit) gates
 3. Measurement: collapsing the superposition of the qubits into either $|0\rangle$ or $|1\rangle$ to get a classical output
- Quantum gates can have parameters and this allows us to make parametrised quantum circuits

Chapter 8

The Fourier Transform

One of the most powerful tools mathematics has brought us is the ability to break down a large system into manageable components. This chapter describes a method that has literally transformed our understanding of the natural world, opening up opportunities in engineering and computer science that has delivered great benefit to humanity.

8.1 Predicting the tides: From Newton to D-Day

As the earth spins on its axis, the gravitational attraction from the moon pulls ever so slightly more on one side of the planet than the other. At the same time, the centripetal acceleration of the rotation exerts a force on the oceans that is almost perfectly provided for by gravity. The combination of these two effects are what lead to the high and low tides occurring twice per day. In this simplified model, it might be assumed that the tides would occur at the same time each day. In reality, the timings of high and low tide drift from one day to the next.

For millennia, the only way to predict the timings of the tides was to consult previous observations. The earliest printed record going back to China in 1056. When he published the Principia, Isaac Newton introduced his theory of gravitation, providing the first scientific explanation for the motion of the tides. Nearly a century passed before Laplace described a set of equations that could model the slow dynamic response of the oceans to the gravitational attraction of the moon and the sun. The complexity of Laplace's equations meant that nearly another century was required before they could be solved.



Figure 8.1: Kelvin's Harmonic Analyser: A mechanical device built in 1878 to compute Fourier series coefficients for tidal predictions. This machine could analyze up to 11 harmonic components simultaneously.

William Thompson, better known later as Lord Kelvin, developed Laplace's approach using the harmonic method. Kelvin's insight was to treat the complex solution as a series of oscillations that could be put together to form the complex

solution. Harmonics, like how strings on a guitar can be plucked together to form a chord, consist of frequencies where each frequency carries some energy. The amplitudes and phases, like the notes and timings of a song, determine how the tidal behaviour can be predicted. Whilst this method made the problem solvable, it was still beyond what could be calculated by hand. To solve this problem, an analogue computer was designed and built.

A far cry from the digital electronic computers of today, the brass machines of the 19th century was built on gears and spheres. A network of pulleys and wires connected the spheres to a pen that would graph out the solution onto paper. The 11 gears in turn correspond to 5 sets of 2 with a final gear corresponding to the mean water level. The speeds at which the shafts rotated was proportional to the frequencies of these physical effects.

8.2 The Fast Fourier Transform

Soon after the second world war ended, the cold war began. Both digital computing and nuclear weapons would transform the nature of warfare. Like how nuclear weapons redefined the unit of explosive power to kilo and megatons of TNT, digital computers would redefine the information processing speed from manual or analogue calculations.

When a measurement is made, some data can be extracted from the measurement device. The data is typically a series of points, separated in increments of space and time. As the separation between each increment gets smaller, the resolution of this discrete dataset improves.

Digital computers are fundamentally limited in the precision they can achieve. Floating point numbers carry a finite number of bits. 32 bit for single precision, 64 for double precision. What this means is that every number stored on every computer is stored to a finite number of decimal points. Even for analogue computers, there is some arbitrary degree to which the precision of a given value can be stored.

8.3 A quantum mechanical treatment

The phase on a qubit can be thought of like the time that has passed during some physical interaction. Like the time read by a clock, this phase contains some information about what happened to the qubit. Just like the time of day is defined by the number of seconds since midnight, the phase represents the total "time" that has passed for a qubit since the quantum circuit was started. When all the qubits are measured, the arrow on the Bloch sphere is projected onto either the top or bottom of the Bloch sphere. In the computational basis, the phase angle has no effect on the measurement angle. Whenever a qubit is measured the probabilities of measuring $|0\rangle$ or $|1\rangle$ are completely independent of whatever phase the qubit might have. When the measurement projects the qubit into either outcome, the phase is effectively destroyed. If the phase the qubit has right before it is measured doesn't change the probabilities, why does it matter?

In chapter 5.3 the coin flip operator was able to distinguish between $|+\rangle$, $|-\rangle$ by a second application. The flip operator for single qubit gates is better referred to as a Hadamard gate. For multiple qubits, the quantum Fourier transform implements this "flipping" where the phase on the qubit can be distinguished from later measurements in the computational basis. When we add more qubits, our computational basis (also known as the dimension of the Hilbert space) expands exponentially. Any measurement outcome of multiple qubits in the computational basis is returned as a bitstring. Converting this bitstring to decimal gives a more readable number, the discrete set of $0, 1, 2, \dots, 2^n - 1$ makes up the computational basis. Unlike the number of seconds on a classical clock (whether it has "hands" or a digital display) the phase on a qubit is analogue (so far as "time" itself is continuous). As with any analogue quantity, we can approximate phase as a binary number. For n bits of precision, we can discretise phase by breaking down the equator of the Bloch sphere into a series of $N = 2^n$ equally spaced points.

Each point in Figure 8.2 corresponds to a specific phase. The angles chosen around the unit circle have a specific pattern. For n qubits, each circle on 8.2 has 2^n points given by $\phi_n^0, \phi_n^1, \dots, \phi_n^{N-1}$

These angles can easily be computed as

$$\phi_n^j = \frac{2\pi j}{2^n}$$

. Starting from the horizontal line on the n^{th} circle, $j = 0$ is the first point on the right side. $j = 1$ is reached by going counter-clockwise to the next point. Then going along each rotation of $j + 1$ until the last point before the first point at $j = N - 1$.

Each point on the circle is separated from its neighbour by $\Delta\phi = \phi_n^{j+1} - \phi_n^j = \frac{2\pi}{N}$. To simplify the maths further on, the convention of using ω to represent the exponential of this distance between the phase points is given as

$$\omega = e^{-i\Delta\phi} = e^{-i2\pi/N}$$

In the computational basis, all these phases result in the same measurement outcomes. By transforming to the Fourier basis, the phases become distinguishable. These basis states, $\{|\tilde{j}\rangle\}$ can also be expressed as a tensor product as

$$\{|\tilde{j}\rangle\} = \frac{1}{\sqrt{N}} \bigotimes_{k=0}^{N-1} (|0\rangle + \omega^{jk} |1\rangle)$$




Figure 8.2: Visualization of the phases that can be represented using $\hat{F}_n$. For n qubits, the Fourier basis contains a discrete set of 2^n phases (points) represented on the complex plane.

With the Fourier basis defined, the operator which translates between the Z basis and the Fourier basis is defined as the quantum Fourier transform

$$\hat{F} |j\rangle = |k\rangle$$

$$\hat{F}_n = \frac{1}{\sqrt{2^n}} \sum_{j,k=0}^{2^n-1} e^{i2\pi jk/2^n} |\tilde{j}\rangle \langle j|$$

Where j, k are any of the number states in the computational (Z) basis. The inverse QFT acts the opposite way, such that any state in the Fourier basis is remapped back to the computational basis. More often than the regular QFT, it is the inverse QFT that is used, as states are typically prepared and measured in the Z basis. The circuit that implements $\hat{F}$ is as




Figure 8.3: Quantum circuit implementing the (inverse) QFT for 5 qubits. Each qubit undergoes a sequence of Hadamard and controlled phase gates. Not shown are the final swap gates needed to reverse the qubit order.

8.4 A new limit of machine precision: Quantum Phase Estimation

In our discrete Fourier transform, ω is quantised, so there is a limit to which the phase of any qubit can be measured to. Therefore, for some arbitrary phase angle ϕ , and for a given number of qubits, there is a limit to the accuracy this phase

can be estimated. Using the QFT, the maximum error to which this phase can be measured is given by the angular distance Δ

$$\Delta = |\phi - 2\pi jk/N|$$

In the best case, $\phi = 0$, so simply picking $j = k = 0$ allows $\Delta = |0 - 0| = 0$. In that case, there is no phase and therefore no error ($\Delta_{min} = 0$). In the worst case, ϕ_{worst} falls between two of the roots of unity: $\phi_{worst} = 2\pi jk/2N$. This maximises the angular distance Δ

$$\Delta_{max} = |2\pi jk/2N - 2\pi jk/N|$$

Since we can always rotate the Bloch sphere by some integer number of $2\pi/N$ without changing Δ , we can simplify this to take $j = k + 1$

$$\Delta_{max} = 2\pi \left| \frac{1}{2N} - \frac{1}{N} \right|$$

Leaving us with the worst case error

$$\Delta_{max} = \frac{\pi}{N} = \frac{\Delta\phi}{2}$$

Assuming an ideal FSFT quantum computer, the (phase) limit of machine precision can be given by:

$$\lim_{n \rightarrow \infty} \Delta_{max} = \frac{\pi}{2^n} \rightarrow 0$$

Since Δ_{max} only depends on n , it might seem as simple as setting having as many qubits as possible such that $N = 2^n$ goes to infinity and Δ_{max} goes to zero. Alas, there are a few more considerations:

- **Connectivity:** All-to-all connectivity is required to implement $\hat{F}$. This can be overcome with SWAP gates at some overhead that scales with n depending on local (logical) qubit connectivity.
- **SPAM:** State Preparation And Measurement. The quantum circuit has to start from the fiduciary state $|0\rangle^{\otimes n}$, prepare the phase corresponding to ϕ_{worst} , carry out QPE and then perform a measurement in the computational basis. Any error in preparing the state or measuring it will in turn raise the error beyond Δ_{max} . Whilst repeated measurements can reduce shot noise, SPAM fidelity will still be a limiting factor, especially if each shot takes a long time to compute.
- **Phase control:** Controlled R_z gates with at least $2\pi/N$ precision. Effectively this defines a limit on the maximum phase noise that can be tolerated by 2-qubit gates. Practically speaking, this sets stringent requirements on the fidelity of the Hadamard gates and the controlled Z rotations. In practice, this fidelity requirement demands effective quantum error correcting codes.
- **Coherence times:** For any given n , F requires $O(n^3)$ circuit layers. Thus the minimum depth of the circuit also scales with $O(n^3)$. Each layer of the circuit requires some time to be processed. Thus it is necessary to maintain a coherent phase for at least the time taken to complete the QPE algorithm.
- **Gate speed:** QPE needs to act much faster than the qubit can dephase. The single qubit and entangling operations that comprise QPE must occur much, much faster than the phase of the qubit can change (say by ϕ_{worst}).

8.5 Faster QFTs

In almost all problems, the set of useful results is much smaller than the total set of possible results. There are many ways I can fit my clothes into my wardrobe, but very few where my shirts don't end up looking like an accordion when some poor soul opens the doors. Often, knowledge of the problem can provide a set of shortcuts. Instead of brute forcing every possible way to fit a shirt in a box, using a hanger greatly reduces the number of possible configurations. Where the quantum computers today have each qubit as a short-lived, expensive and noisy resource, clever tricks can be made to squeeze out every (qu)bit of performance. The considerations outlined in the previous section pose some daunting engineering challenges before the QFT can be of any practical use to anyone. As the FFT enabled 20th century computers to vastly expand their problem solving capability, perhaps faster QFTs can accelerate the adoption of quantum computing in the 21st century.

8.6 Additional resources

- Qiskit Lab 5: Accuracy of Quantum Phase Estimation
- 3blue1brown But What is the Fourier Transform? A visual introduction